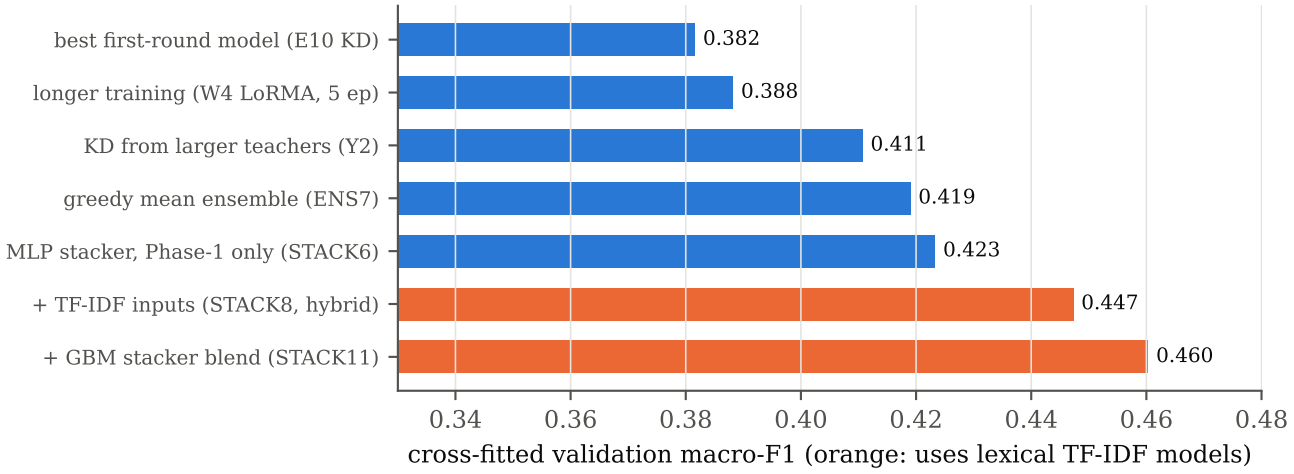
best first-round model (E10 KD)
0.382
longer training (W4 LoRMA, 5 ep)
0.388
KD from larger teachers (Y2)
0.411
greedy mean ensemble (ENS7)
0.419
MLP stacker, Phase-1 only (STACK6)
0.423
+ TF-IDF inputs (STACK8, hybrid)
0.447
+ GBM stacker blend (STACK11)
0.460
0.34
0.36
0.38
0.40
0.42
0.44
0.46
0.48
cross-fitted validation macro-F1 (orange: uses lexical TF-IDF models)